

✓ AI & Machine Learning Internship Projects | Future Interns

Welcome to this Google Colab workspace! This notebook serves as a comprehensive portfolio showcasing my work during the **Future Interns Machine Learning Internship** program.

Throughout this internship, I'm developing practical machine learning solutions for real-world business challenges. This file will contain detailed explanations, code implementations, and visual insights for various tasks undertaken.

Currently, this notebook focuses on:

- **Task 2: Churn Prediction System** - Building a machine learning model to identify customers at risk of churning, providing valuable insights for retention strategies.

You'll find a structured approach, from data loading and exploration to model building, evaluation, and deriving actionable business recommendations.

Let's explore the power of Machine Learning in solving business problems!

1. Project Setup & Data Loading

First, we'll set up our environment by importing necessary libraries and loading the `WA_Fn-UseC_-Telco-Customer-Churn.csv` dataset. This dataset contains various customer attributes and their churn status.

The 'TotalCharges' column in this dataset often loads as a string due to some empty values, which we'll correct by converting it to a numeric type and handling the resulting missing values. We will also convert the 'Churn' target variable into a numerical format (0 for 'No', 1 for 'Yes').

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, confusion_matrix, roc_curve, auc

# --- Data Loading ---

# Since you've uploaded the file directly, it will be in the Colab session's root directory.
# The 'WA_Fn-UseC_-Telco-Customer-Churn.csv' is the file name.
file_name = 'WA_Fn-UseC_-Telco-Customer-Churn.csv'

# Check if the file is already in the current directory (after initial upload)
import os
if not os.path.exists(file_name):
    print(f"File '{file_name}' not found. Please upload your customer churn dataset.")
    from google.colab import files
    uploaded = files.upload()
    # If the user uploads with a different name, adjust file_name
    if uploaded and list(uploaded.keys())[0] != file_name:
        print(f"Warning: Uploaded file is named '{list(uploaded.keys())[0]}'. Using this name.")
        file_name = list(uploaded.keys())[0]

df = pd.read_csv(file_name)

# Display basic information about the dataset
print("--- Dataset Head ---")
print(df.head())

print("\n--- Dataset Info ---")
df.info()

print("\n--- Missing Values ---")
print(df.isnull().sum())


# --- Important Initial Preprocessing for Telco Churn Dataset ---
# The 'TotalCharges' column in this dataset often gets loaded as an object (string)
# and contains some empty strings which cause issues when converting to numeric.
# We'll convert it to numeric and handle potential missing values (coerce errors to NaN).
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')

# Handle rows where 'TotalCharges' became NaN (typically empty strings representing new customers)
# For this dataset, dropping these rows is usually acceptable as they are very few.
# This ensures all values in 'TotalCharges' are numerical.
df.dropna(subset=['TotalCharges'], inplace=True)

# Convert the 'Churn' target variable to numerical (0 and 1) for modeling
# 'Yes' usually indicates churn (1), 'No' indicates no churn (0).
df['Churn'] = df['Churn'].apply(lambda x: 1 if x == 'Yes' else 0)

print("\n--- Processed DataFrame Head (after TotalCharges and Churn conversion) ---")
print(df.head())

print("\n--- Churn Column Value Counts (after conversion) ---")
print(df['Churn'].value_counts())
```

 --- Dataset Head ---

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	7590-VHVEG	Female	0	Yes	No	1	No	
1	5575-GNVDE	Male	0	No	No	34	Yes	
2	3668-QPYBK	Male	0	No	No	2	Yes	
3	7795-CFOCW	Male	0	No	No	45	No	

```
4  9237-HQITU  Female          0      No      No      2      Yes

      MultipleLines  InternetService  OnlineSecurity  ...  DeviceProtection  \
0  No phone service          DSL          No  ...          No
1          No          DSL          Yes  ...          Yes
2          No          DSL          Yes  ...          No
3  No phone service          DSL          Yes  ...          Yes
4          No      Fiber optic          No  ...          No

      TechSupport  StreamingTV  StreamingMovies      Contract  PaperlessBilling  \
0          No          No          No  Month-to-month          Yes
1          No          No          No      One year          No
2          No          No          No  Month-to-month          Yes
3          Yes          No          No      One year          No
4          No          No          No  Month-to-month          Yes

      PaymentMethod  MonthlyCharges  TotalCharges  Churn
0      Electronic check          29.85          29.85    No
1          Mailed check          56.95          1889.5    No
2          Mailed check          53.85          108.15   Yes
3  Bank transfer (automatic)          42.30          1840.75    No
4      Electronic check          70.70          151.65   Yes
```

[5 rows x 21 columns]

```
--- Dataset Info ---
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column              Non-Null Count  Dtype
---  -
0   customerID          7043 non-null  object
1   gender              7043 non-null  object
2   SeniorCitizen       7043 non-null  int64
3   Partner             7043 non-null  object
4   Dependents          7043 non-null  object
5   tenure              7043 non-null  int64
6   PhoneService        7043 non-null  object
7   MultipleLines       7043 non-null  object
8   InternetService     7043 non-null  object
9   OnlineSecurity      7043 non-null  object
10  OnlineBackup        7043 non-null  object
11  DeviceProtection    7043 non-null  object
12  TechSupport         7043 non-null  object
13  StreamingTV         7043 non-null  object
14  StreamingMovies     7043 non-null  object
15  Contract            7043 non-null  object
16  PaperlessBilling    7043 non-null  object
17  PaymentMethod       7043 non-null  object
18  MonthlyCharges      7043 non-null  float64
19  TotalCharges        7043 non-null  object
20  Churn               7043 non-null  object
```


2. Exploratory Data Analysis (EDA) Understanding the characteristics of our customer data is crucial for building an effective churn prediction model. This section explores the distribution of churn and key features, identifying patterns that might influence churn.

Churn Distribution We'll first examine the balance of our target variable, 'Churn', to understand how many customers churn versus how many do not. This is important as imbalanced datasets can affect model training.

```
# CODE BLOCK 2.1: Churn Distribution

# --- Target Variable Distribution ---
print("\n--- Churn Distribution ---")
print(df['Churn'].value_counts())
print(f"Churn Rate: {df['Churn'].mean():.2%}")

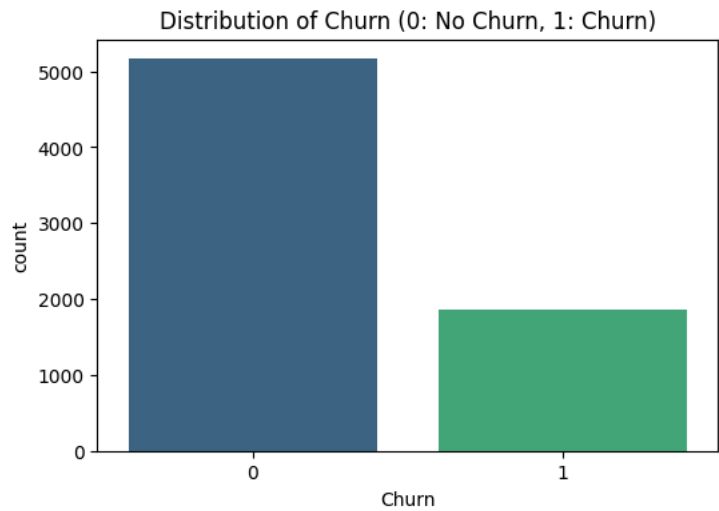
plt.figure(figsize=(6, 4))
sns.countplot(x='Churn', data=df, palette='viridis')
plt.title('Distribution of Churn (0: No Churn, 1: Churn)')
plt.show()
```



```
--- Churn Distribution ---
Churn
0    5163
1    1869
Name: count, dtype: int64
Churn Rate: 26.58%
/tmp/ipython-input-9-1436504502.py:9: FutureWarning:

    Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`

sns.countplot(x='Churn', data=df, palette='viridis')
```



Numerical Feature Distributions Next, we'll visualize the distributions of numerical features like tenure, MonthlyCharges, and TotalCharges. Box plots will help us see how these features differ between churning and non-churning customers.

```
# CODE BLOCK 2.2: Numerical Feature Distributions

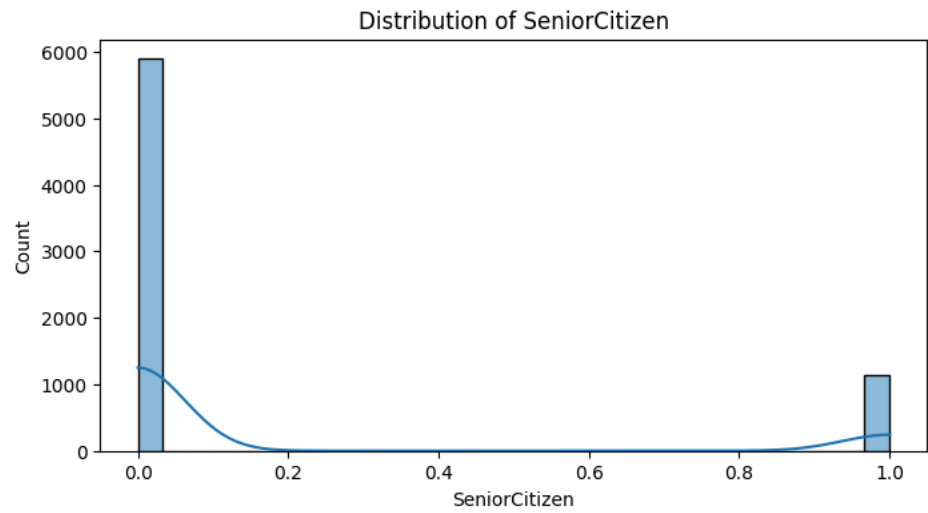
numerical_features = df.select_dtypes(include=np.number).columns.tolist()
# Exclude 'customerID' if it's in numerical_features (already dropped in preproc) and 'Churn' as it's the target
if 'customerID' in numerical_features: numerical_features.remove('customerID') # Guard against ID remaining
if 'Churn' in numerical_features: numerical_features.remove('Churn')

print("\n--- Numerical Feature Distributions ---")
for col in numerical_features:
    plt.figure(figsize=(8, 4))
    sns.histplot(df[col], kde=True, bins=30, palette='viridis')
    plt.title(f'Distribution of {col}')
    plt.show()

plt.figure(figsize=(8, 4))
sns.boxplot(x='Churn', y=col, data=df, palette='viridis')
plt.title(f'{col} vs. Churn')
plt.show()
```

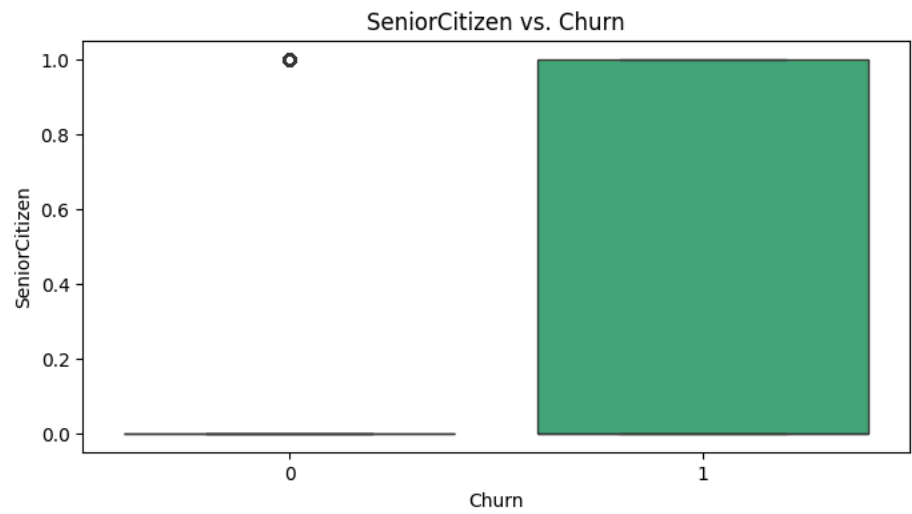


```
--- Numerical Feature Distributions ---
/tmp/ipython-input-10-2647694609.py:11: UserWarning: Ignoring `palette` because no `hue` variable has been assigned.
sns.histplot(df[col], kde=True, bins=30, palette='viridis')
```

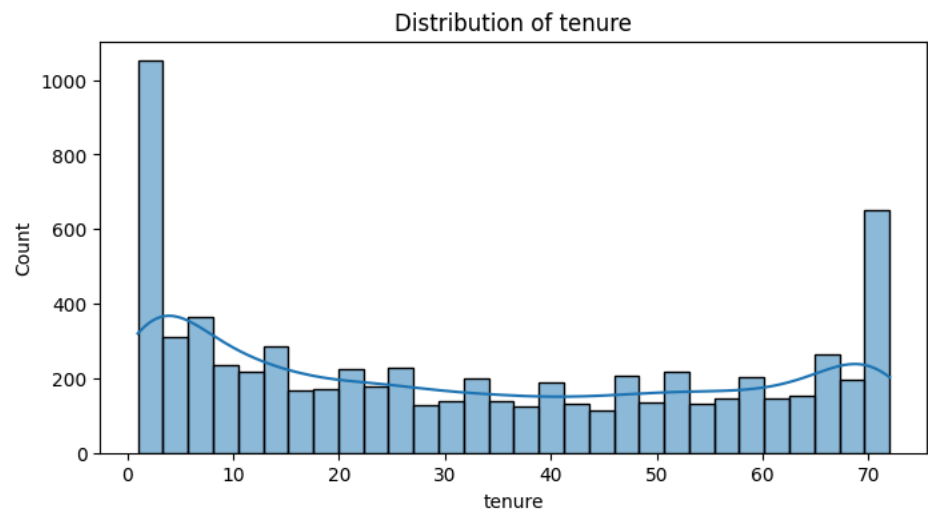


```
/tmp/ipython-input-10-2647694609.py:16: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`

sns.boxplot(x='Churn', y=col, data=df, palette='viridis')
```

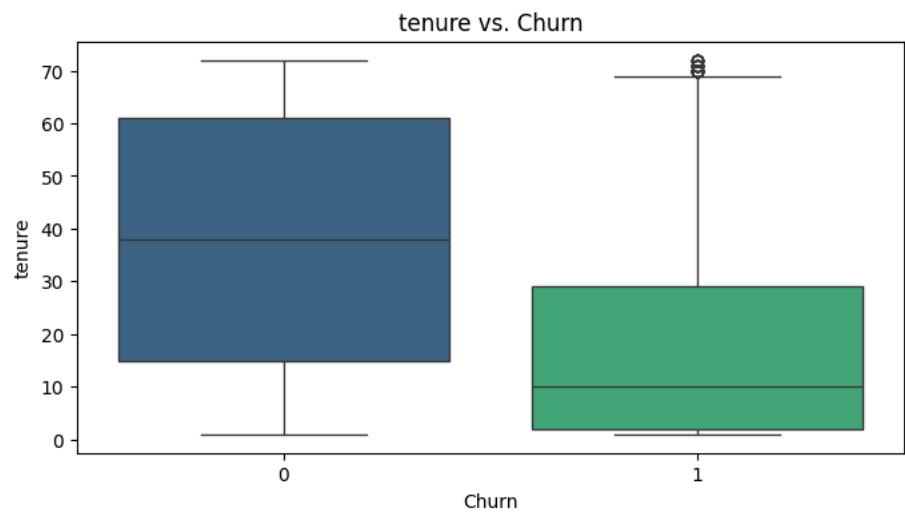


```
/tmp/ipython-input-10-2647694609.py:11: UserWarning: Ignoring `palette` because no `hue` variable has been assigned.
sns.histplot(df[col], kde=True, bins=30, palette='viridis')
```

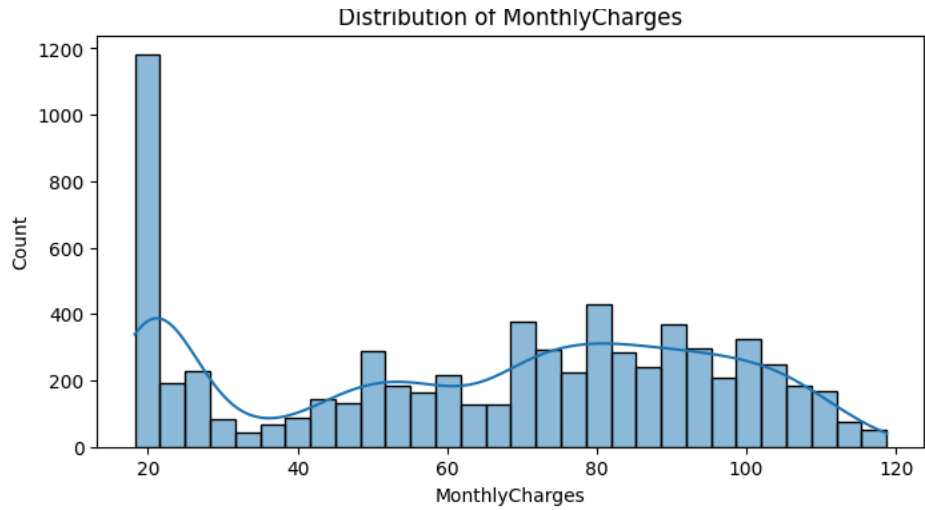


```
/tmp/ipython-input-10-2647694609.py:16: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`

sns.boxplot(x='Churn', y=col, data=df, palette='viridis')
```



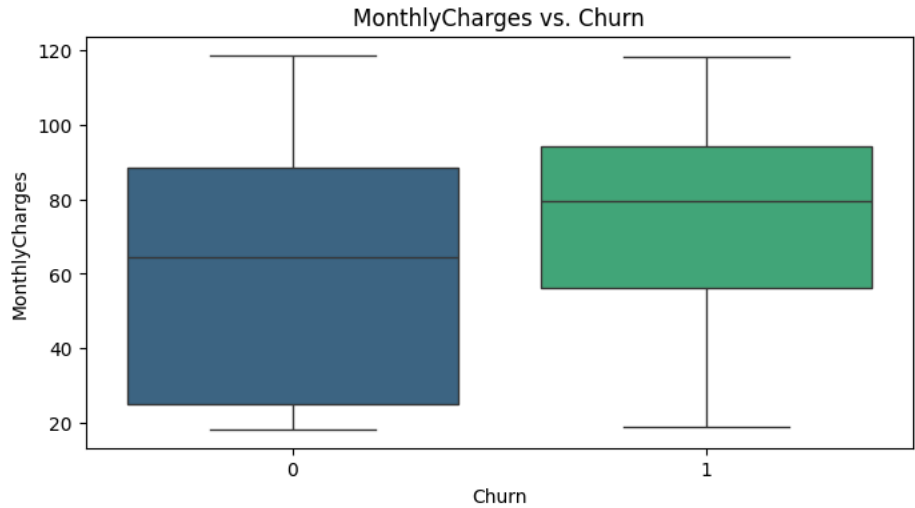
```
/tmp/ipython-input-10-2647694609.py:11: UserWarning: Ignoring `palette` because no `hue` variable has been assigned.
sns.histplot(df[col], kde=True, bins=30, palette='viridis')
```



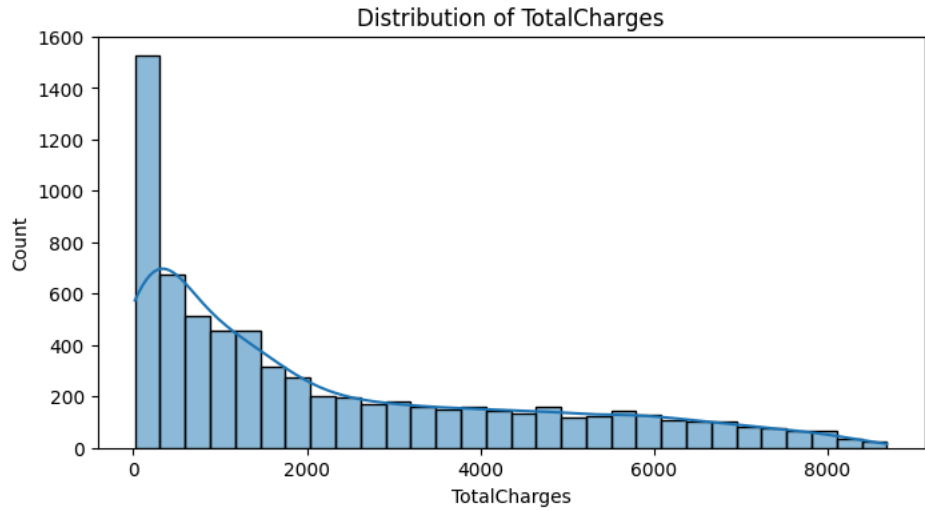
/tmp/ipython-input-10-2647694609.py:16: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`

```
sns.boxplot(x='Churn', y=col, data=df, palette='viridis')
```



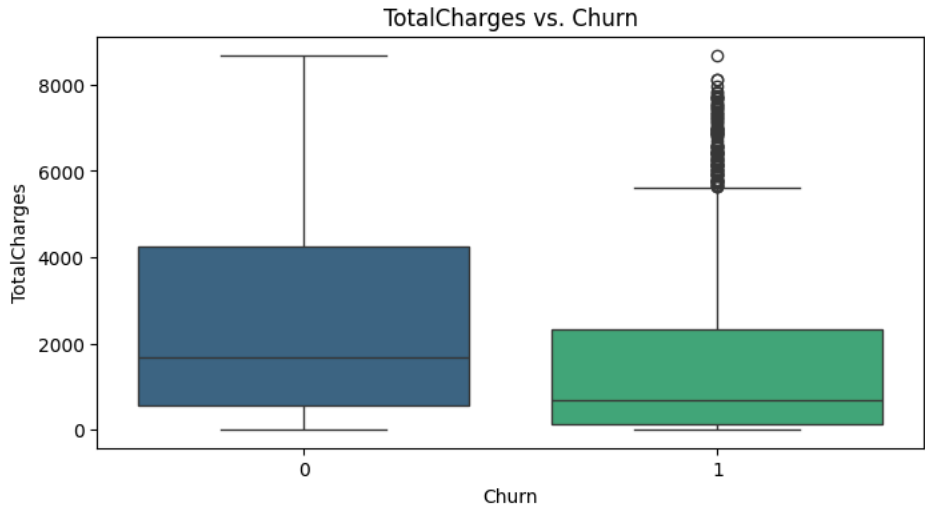
/tmp/ipython-input-10-2647694609.py:11: UserWarning: Ignoring `palette` because no `hue` variable has been assigned.  
sns.histplot(df[col], kde=True, bins=30, palette='viridis')



/tmp/ipython-input-10-2647694609.py:16: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`

```
sns.boxplot(x='Churn', y=col, data=df, palette='viridis')
```



Categorical Feature Churn Rates Finally, we'll explore categorical features like Contract, PaymentMethod, InternetService, etc. by visualizing the churn distribution within each category. This helps identify which categories have higher churn rates.

```
# CODE BLOCK 2.3: Categorical Feature Churn Rates

# Drop 'customerID' early if it wasn't dropped already (e.g. if you reran only EDA)
if 'customerID' in df.columns:
    df_eda = df.drop('customerID', axis=1)
else:
    df_eda = df.copy() # Use the already cleaned df if customerID is gone

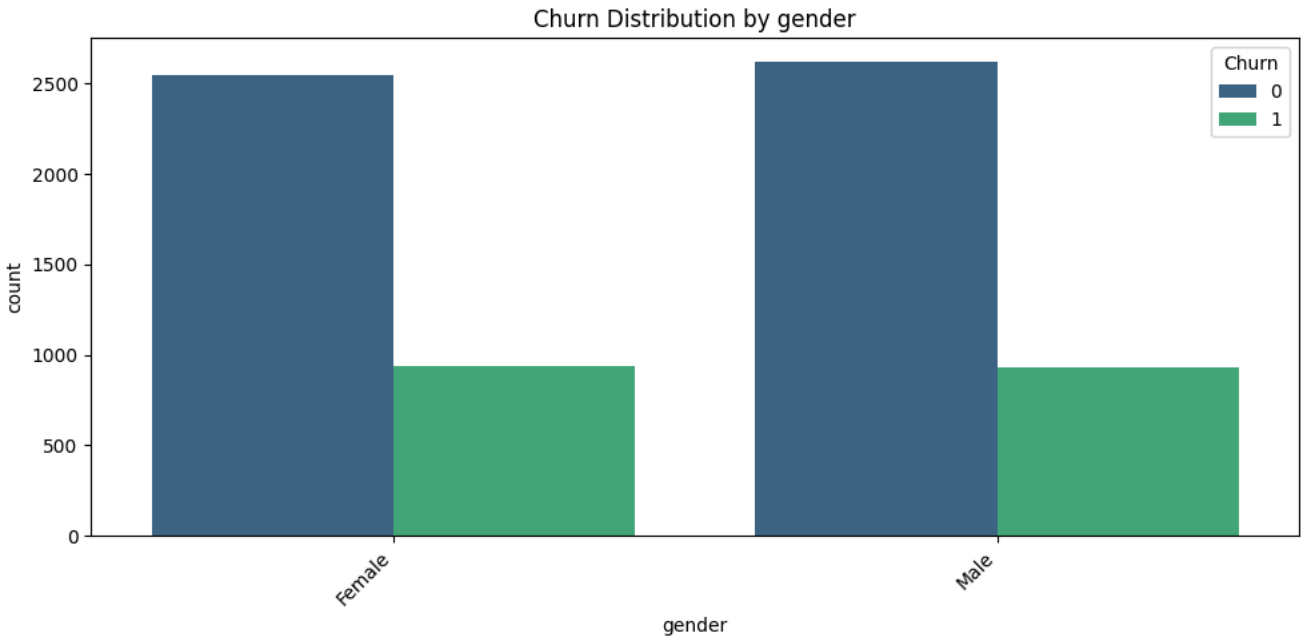
categorical_features = df_eda.select_dtypes(include='object').columns.tolist()

print("\n--- Categorical Feature Churn Rates ---")
for col in categorical_features:
    plt.figure(figsize=(10, 5))
    sns.countplot(x=col, hue='Churn', data=df_eda, palette='viridis')
    plt.title(f'Churn Distribution by {col}')
    plt.xticks(rotation=45, ha='right')
    plt.tight_layout()
    plt.show()

# Calculate and print churn rate per category
churn_rate_by_cat = df_eda.groupby(col)['Churn'].mean().reset_index()
print(f"\nChurn Rate by {col}:\n{churn_rate_by_cat}")
```

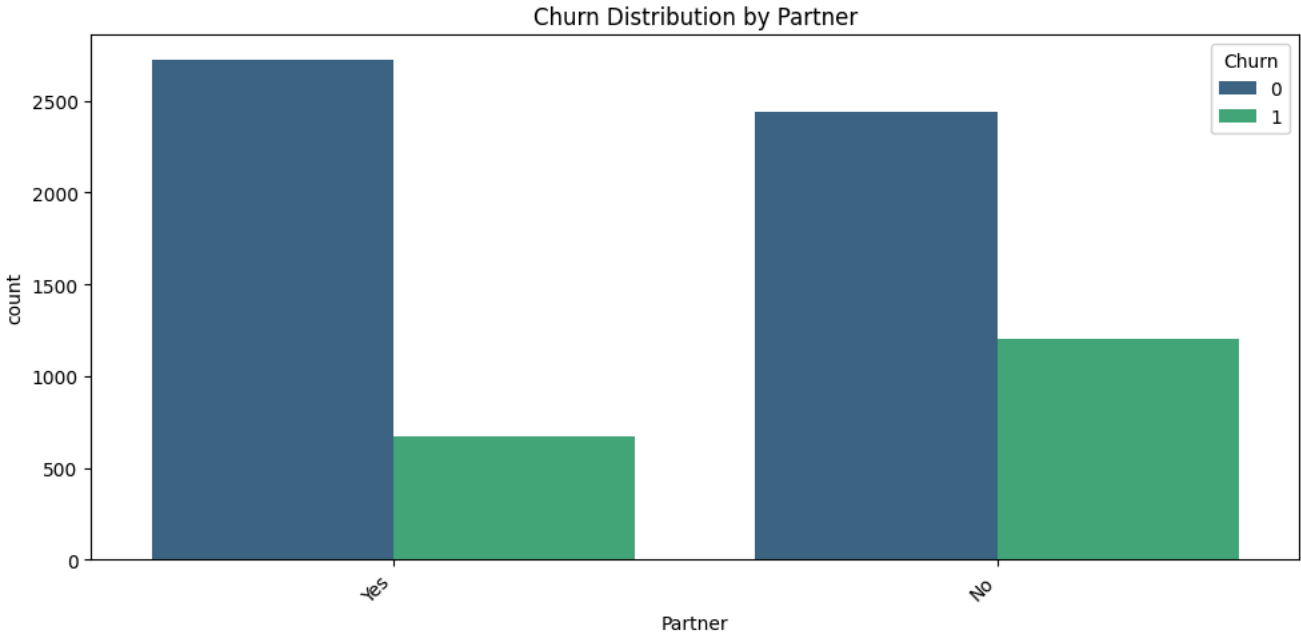


--- Categorical Feature Churn Rates ---



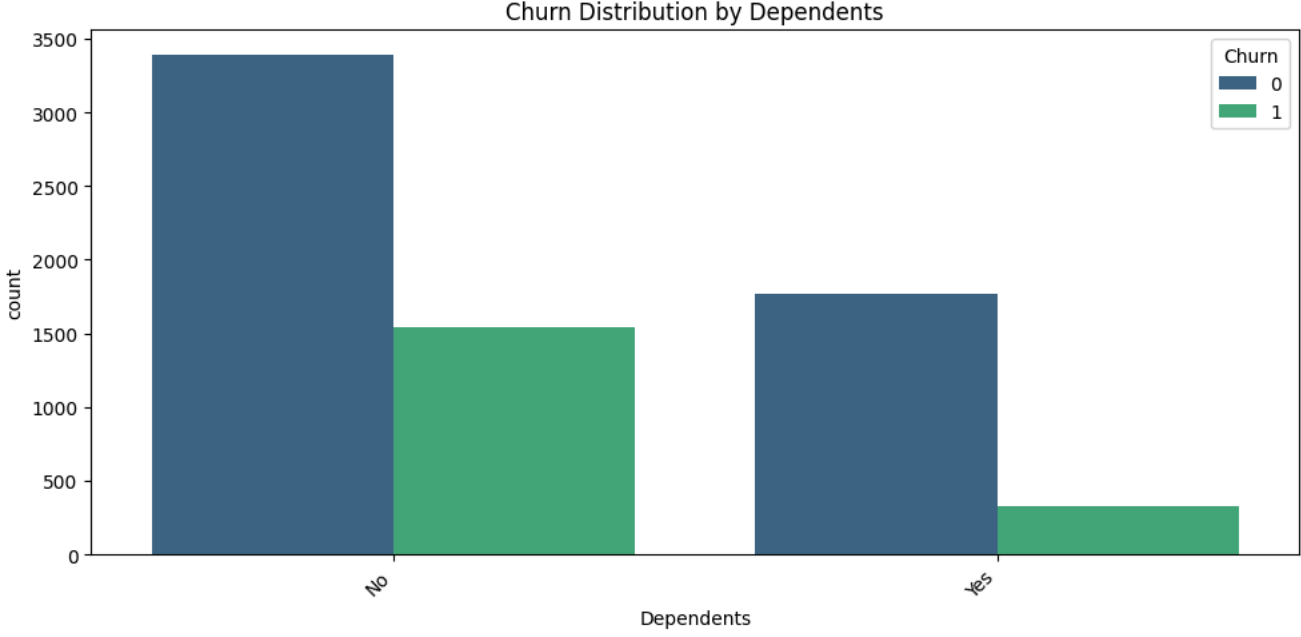
Churn Rate by gender:

gender	Churn
0	Female 0.269595
1	Male 0.262046



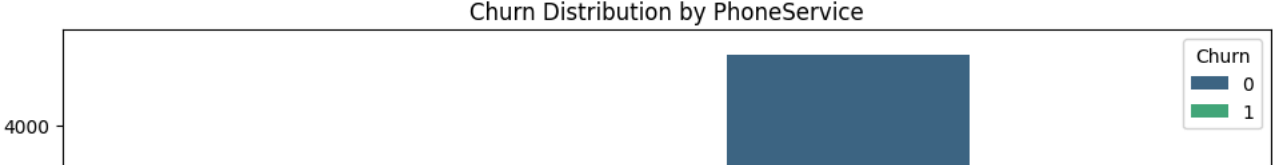
Churn Rate by Partner:

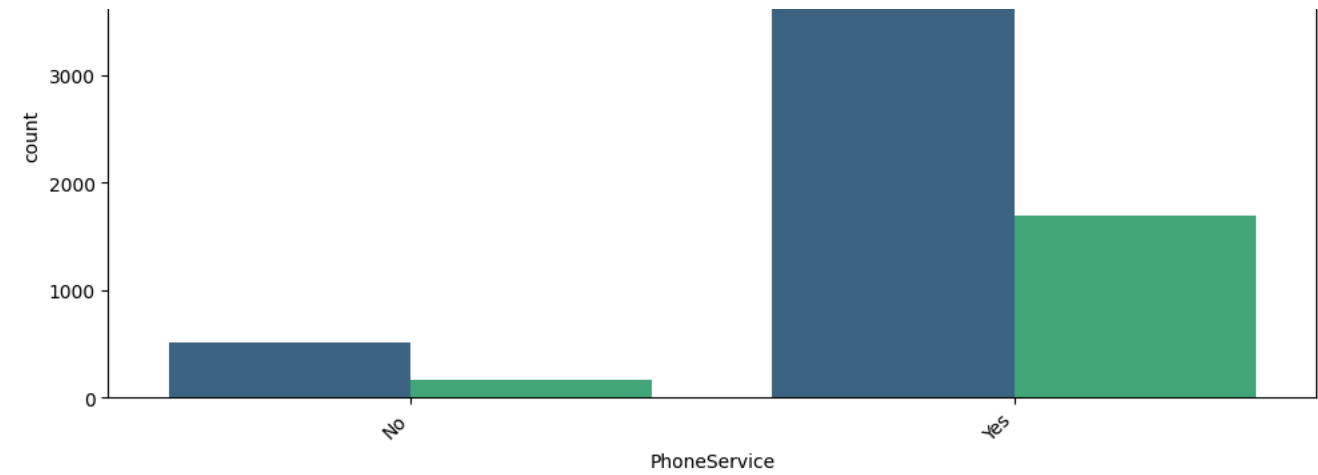
Partner	Churn
0	No 0.329761
1	Yes 0.197171



Churn Rate by Dependents:

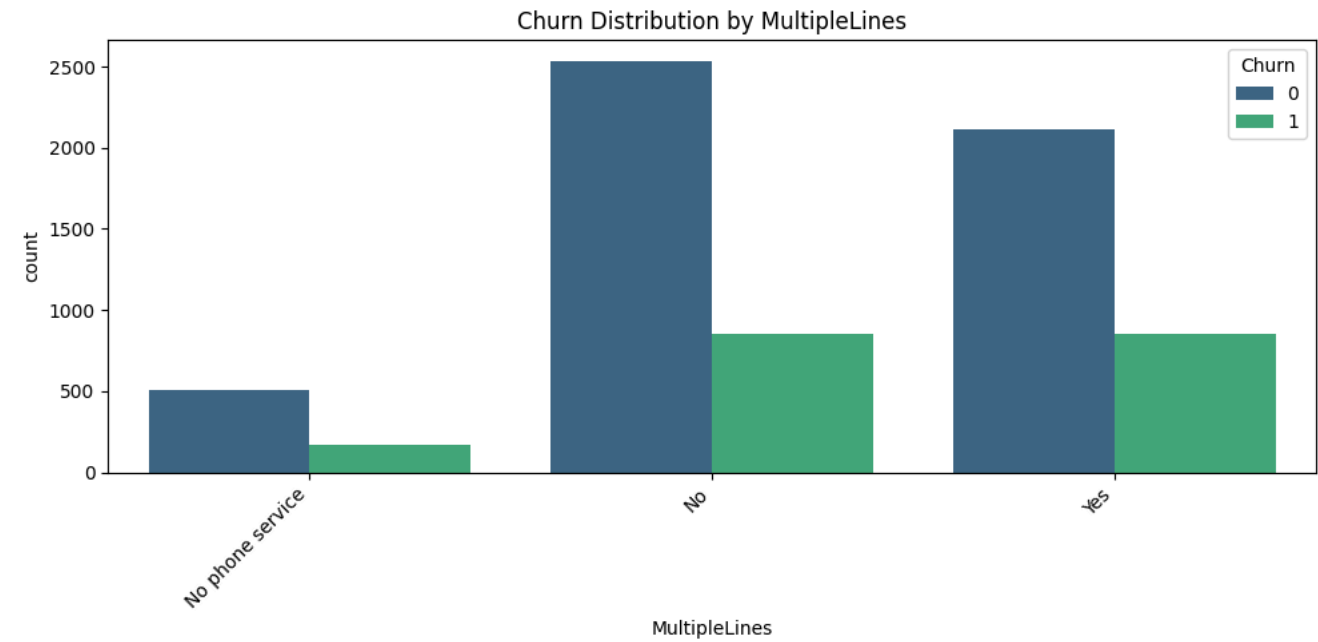
Dependents	Churn
0	No 0.312791
1	Yes 0.155312





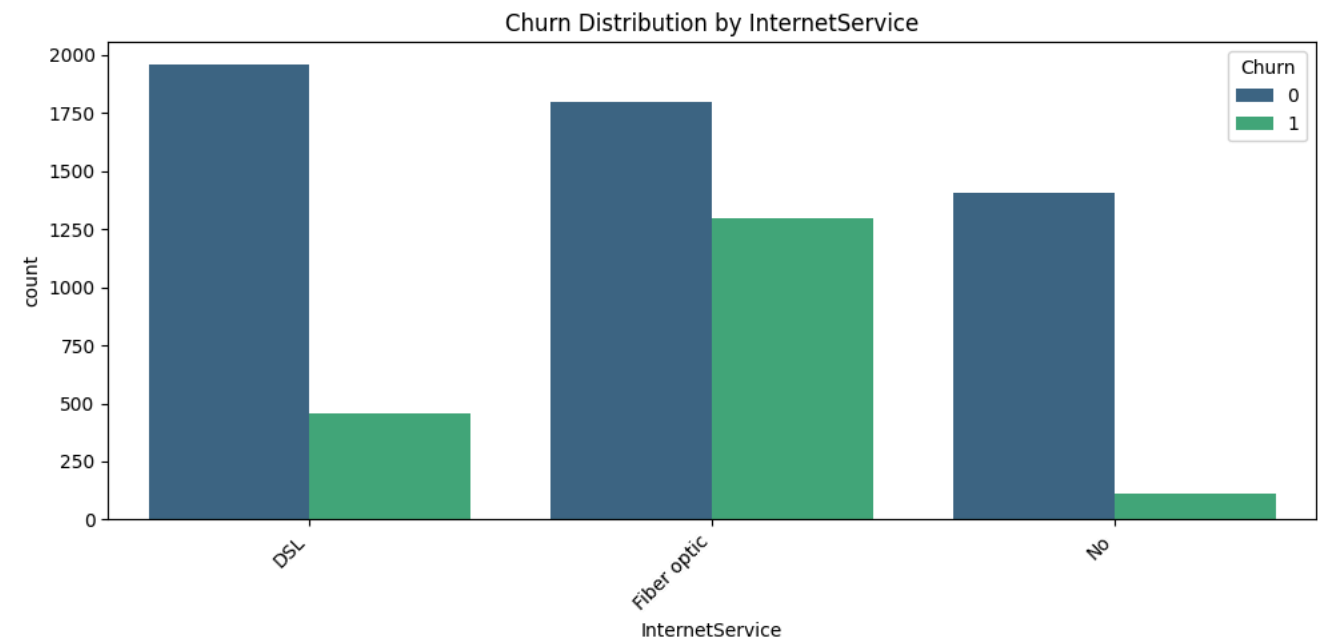
Churn Rate by PhoneService:

PhoneService	Churn
0	No 0.250000
1	Yes 0.267475



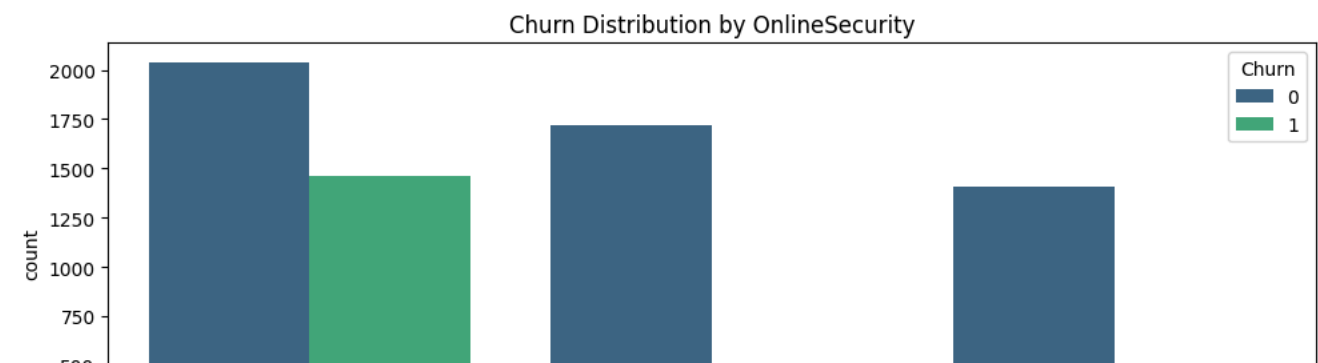
Churn Rate by MultipleLines:

MultipleLines	Churn
0	No 0.250812
1	No phone service 0.250000
2	Yes 0.286485



Churn Rate by InternetService:

InternetService	Churn
0	DSL 0.189983
1	Fiber optic 0.418928
2	No 0.074342

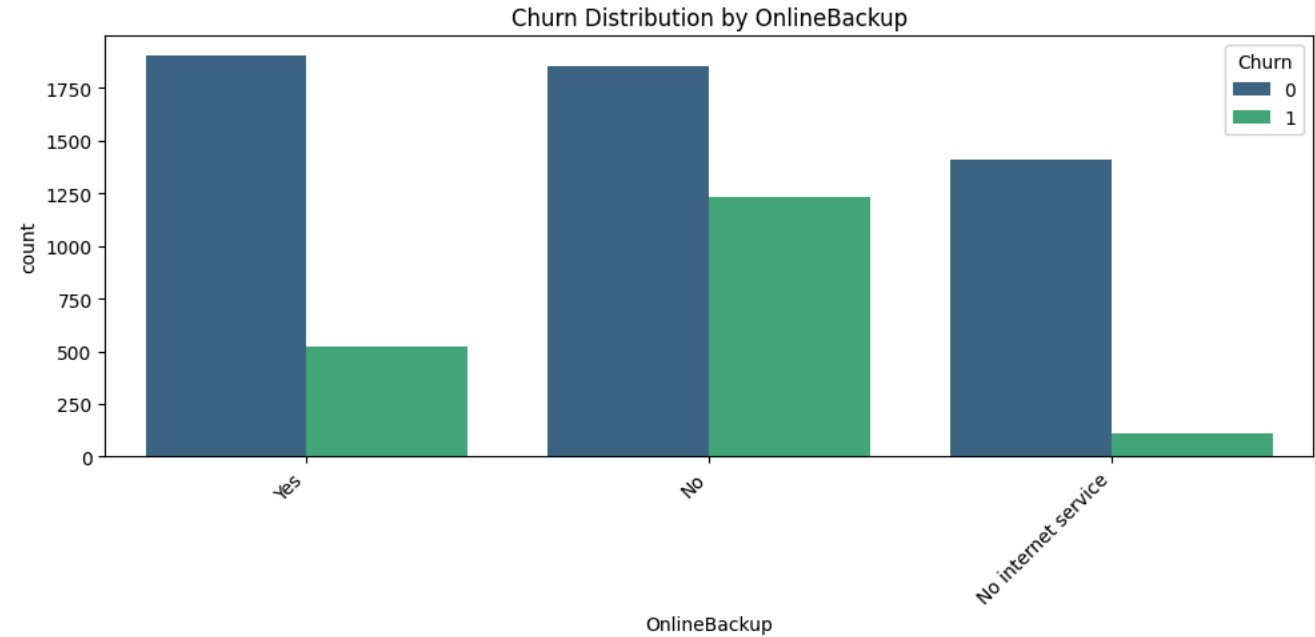






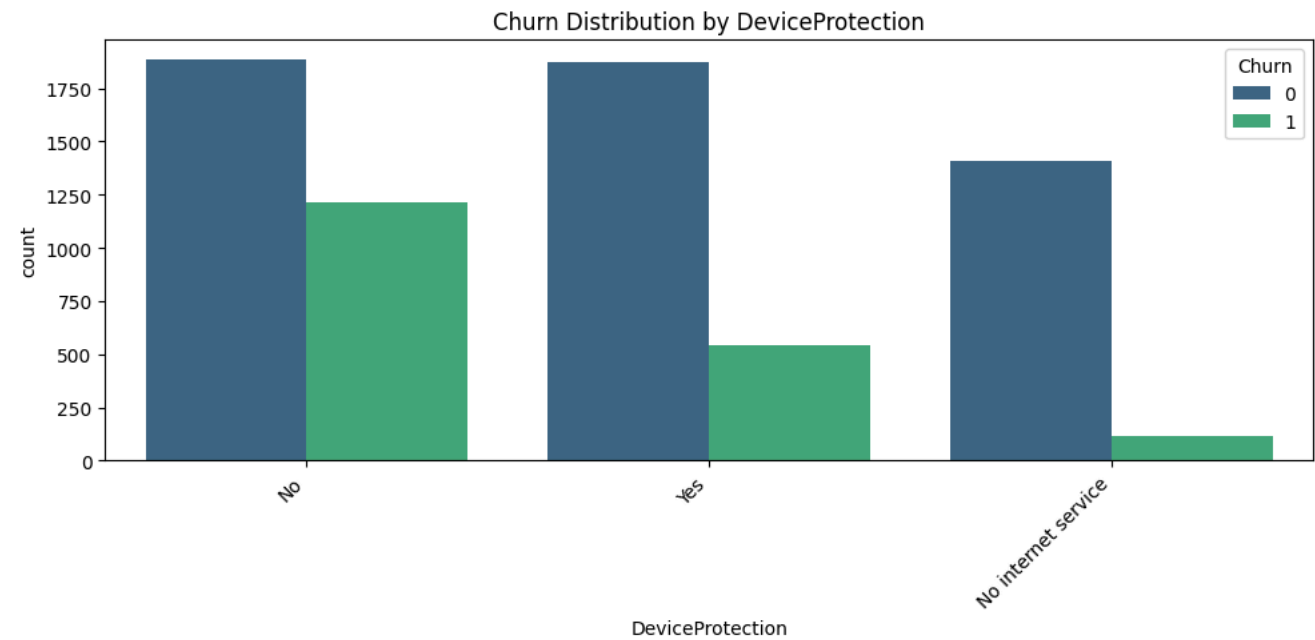
Churn Rate by OnlineSecurity:

	OnlineSecurity	Churn
0	No	0.417787
1	No internet service	0.074342
2	Yes	0.146402



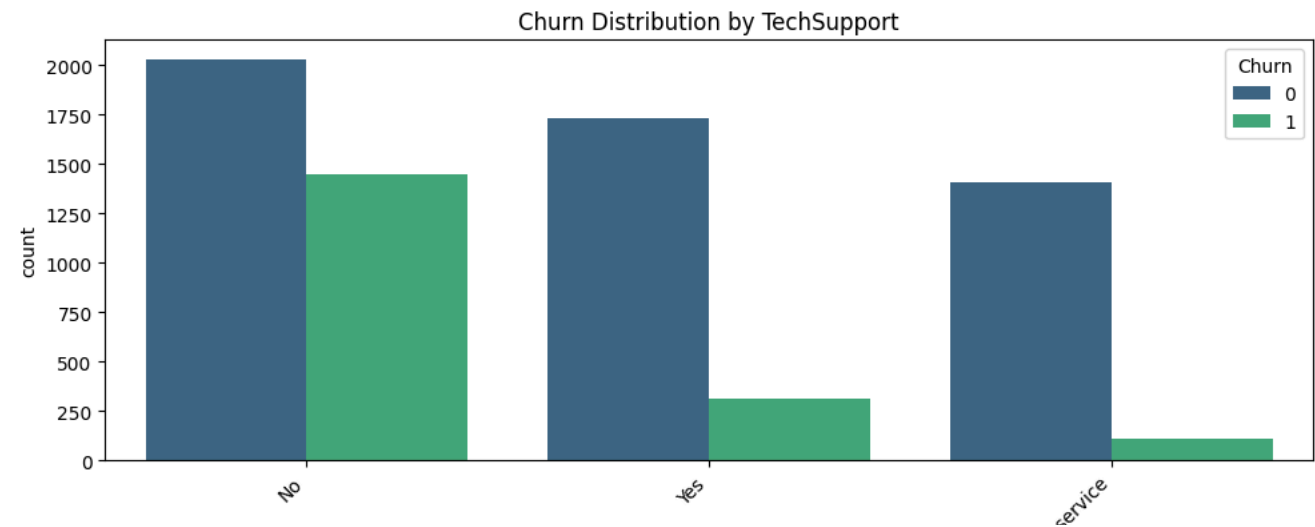
Churn Rate by OnlineBackup:

	OnlineBackup	Churn
0	No	0.399417
1	No internet service	0.074342
2	Yes	0.215670



Churn Rate by DeviceProtection:

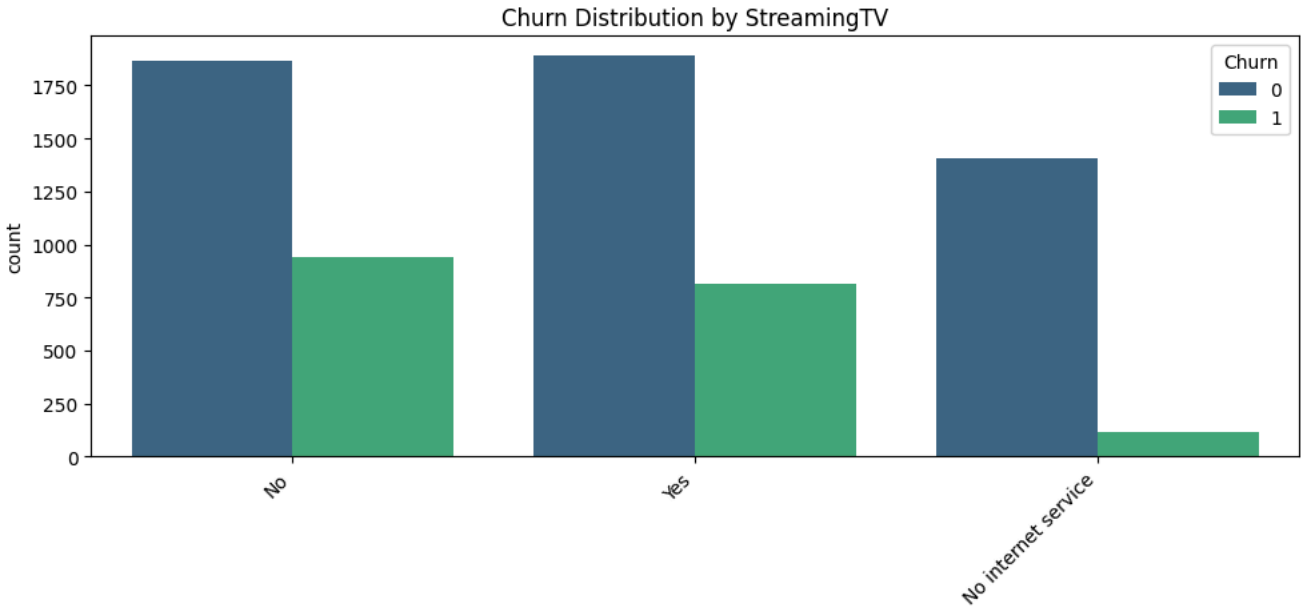
	DeviceProtection	Churn
0	No	0.391403
1	No internet service	0.074342
2	Yes	0.225393



TechSupport

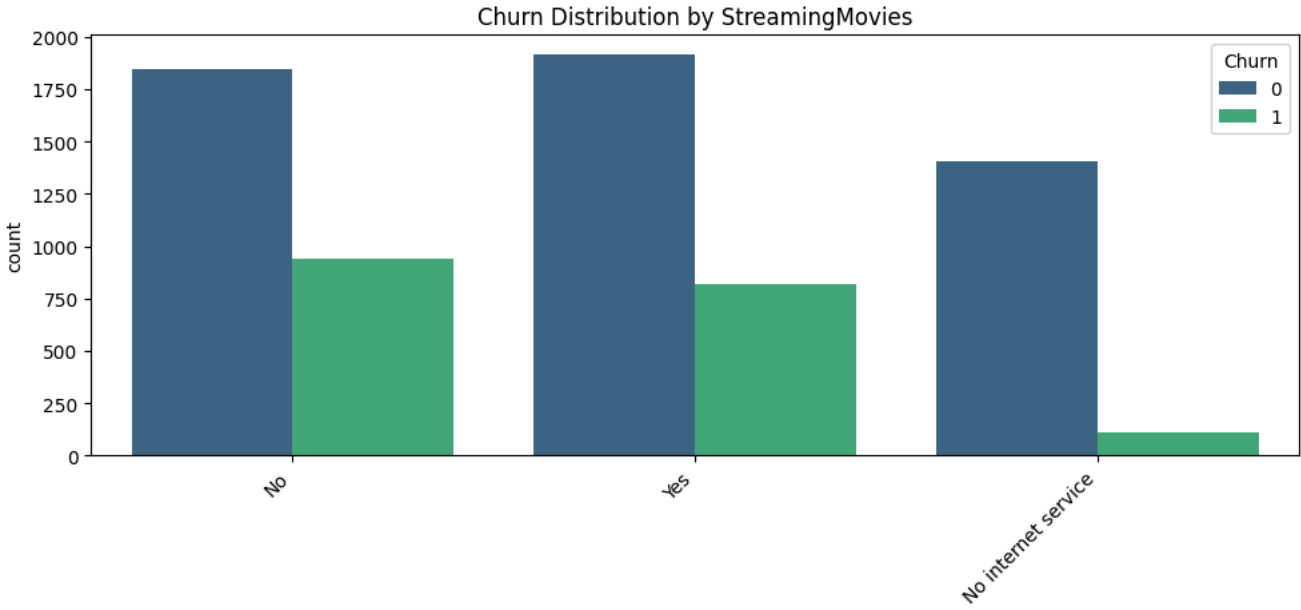
No internet service

Churn Rate by TechSupport:		
	TechSupport	Churn
0	No	0.416475
1	No internet service	0.074342
2	Yes	0.151961



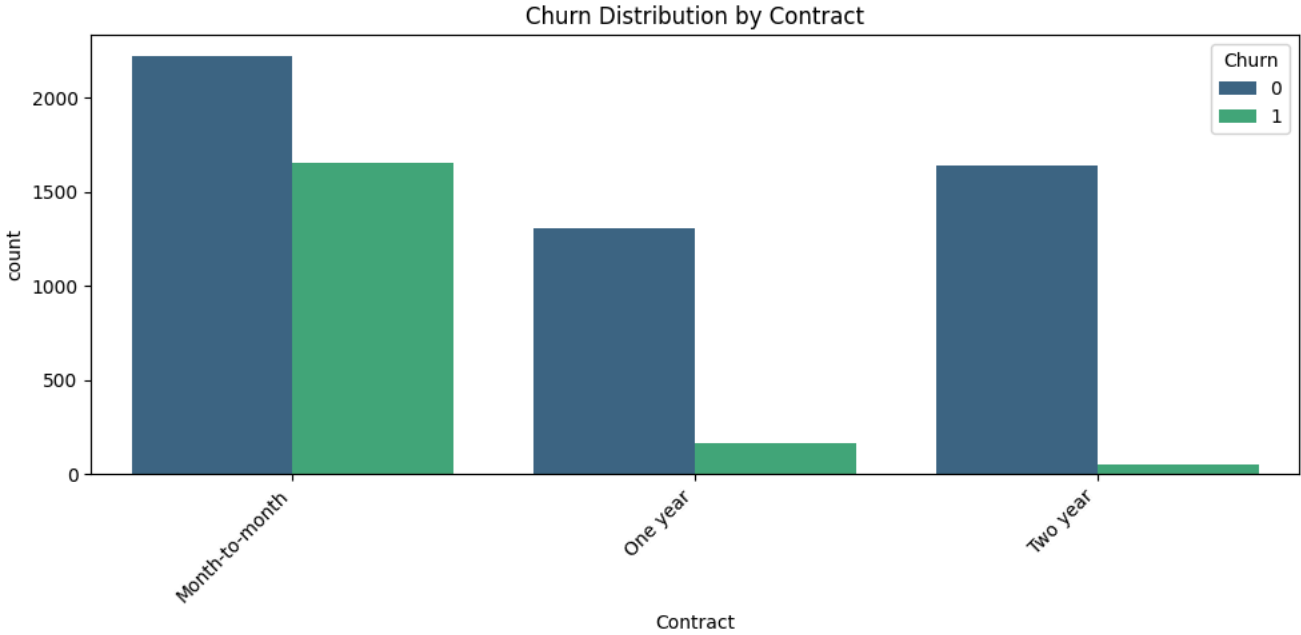
StreamingTV

Churn Rate by StreamingTV:		
	StreamingTV	Churn
0	No	0.335351
1	No internet service	0.074342
2	Yes	0.301147



StreamingMovies

Churn Rate by StreamingMovies:		
	StreamingMovies	Churn
0	No	0.337289
1	No internet service	0.074342
2	Yes	0.299524

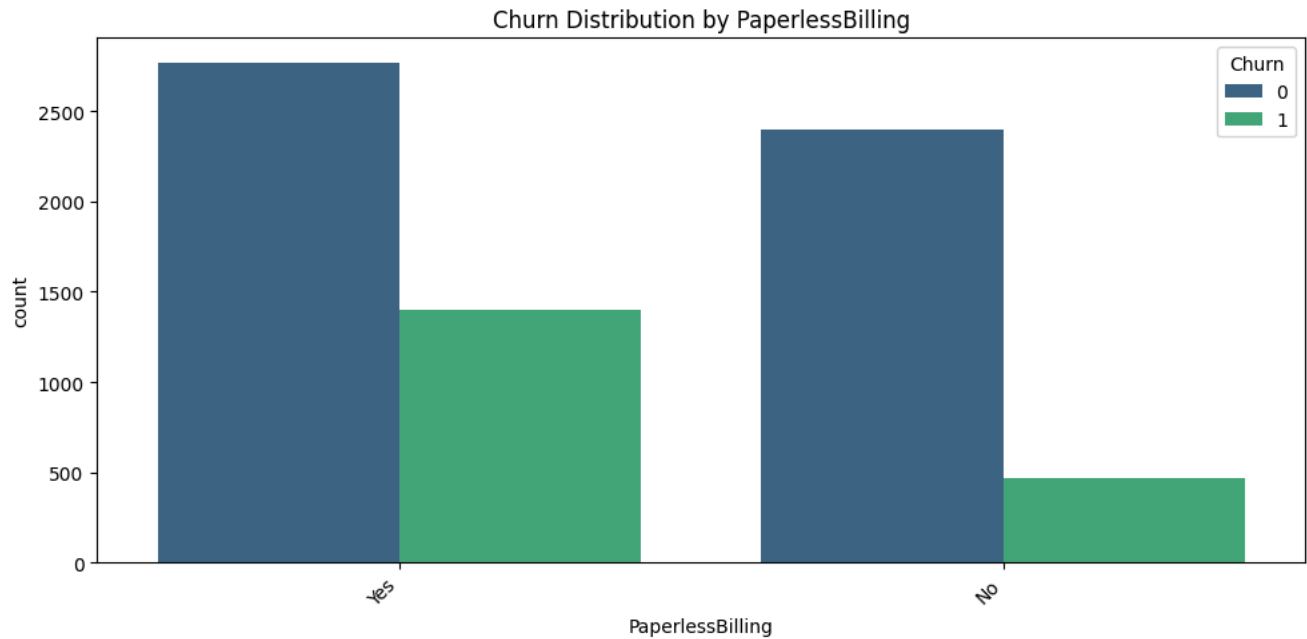


Contract

Churn Rate by Contract:		
-------------------------	--	--

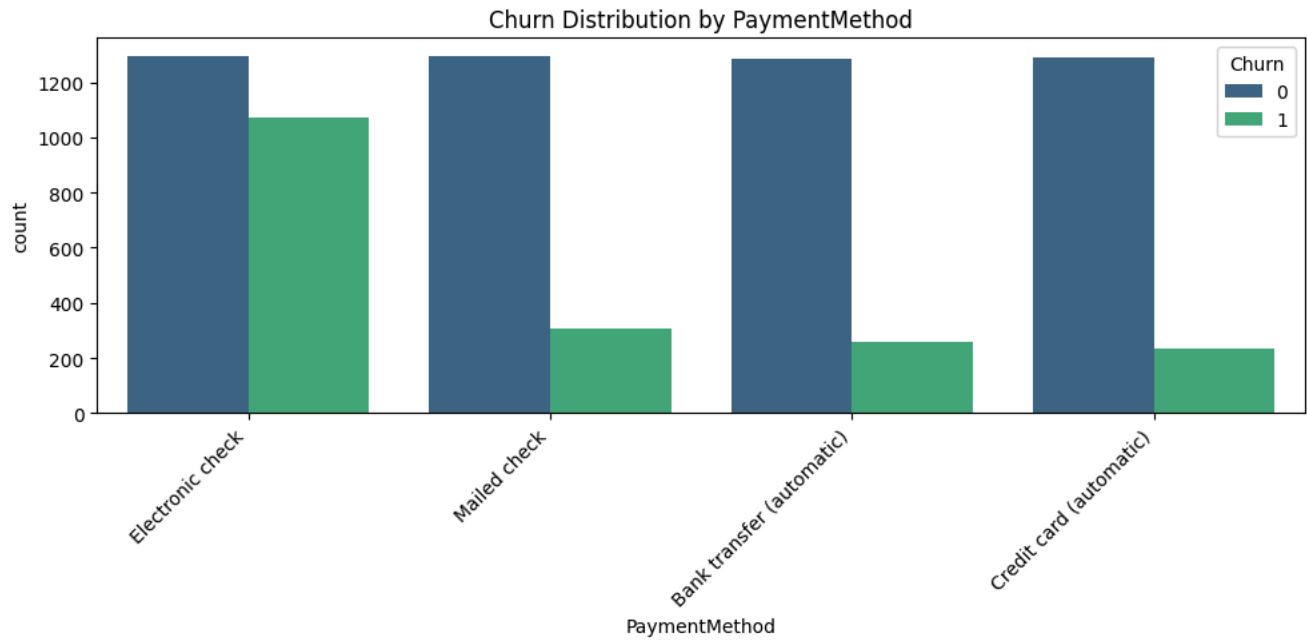
Churn Rate by Contract:

	Contract	Churn
0	Month-to-month	0.427097
1	One year	0.112772
2	Two year	0.028487



Churn Rate by PaperlessBilling:

	PaperlessBilling	Churn
0	No	0.163757
1	Yes	0.335893



Churn Rate by PaymentMethod:

	PaymentMethod	Churn
0	Bank transfer (automatic)	0.167315
1	Credit card (automatic)	0.152531
2	Electronic check	0.452854
3	Mailed check	0.103020



3. Data Preprocessing & Feature Engineering Before training our machine learning models, we need to preprocess the data. This involves encoding categorical variables into a numerical format that models can understand and scaling numerical features to prevent dominance by features with larger values. We will set up a robust preprocessing pipeline using ColumnTransformer for this.

```
# CODE BLOCK 3: Data Preprocessing & Feature Engineering

# Drop 'customerID' as it's just an identifier and not a predictive feature
# We do this here explicitly in case EDA was run on a df with it.
df_for_modeling = df.drop('customerID', axis=1)

# Define target variable and features
X = df_for_modeling.drop('Churn', axis=1)
y = df_for_modeling['Churn']

# Identify categorical and numerical features for preprocessing
numerical_cols = X.select_dtypes(include=np.number).columns.tolist()
categorical_cols = X.select_dtypes(include='object').columns.tolist()

print(f"\nNumerical columns identified: {numerical_cols}")
print(f"Categorical columns identified: {categorical_cols}")

# Create preprocessing pipelines for numerical and categorical features
numerical_transformer = Pipeline(steps=[
    ('scaler', StandardScaler()) # Standardize numerical features
])

categorical_transformer = Pipeline(steps=[
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical features
])

# Create a preprocessor using ColumnTransformer
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numerical_transformer, numerical_cols),
        ('cat', categorical_transformer, categorical_cols)
    ],
    remainder='passthrough' # Keep other columns if any, though none expected here
)

# --- Train-Test Split ---
# We use stratify=y to ensure that the proportion of churned (1) and non-churned (0)
# customers is maintained in both the training and testing sets.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

print(f"\nTraining set shape: {X_train.shape}")
print(f"Testing set shape: {X_test.shape}")

# --- Verify Churn Distribution in Splits (TROUBLESHOOTING STEP) ---
# This helps confirm that your splits contain both classes, preventing the ValueError.
print("\n--- Churn Distribution in Training Set (y_train) ---")
print(y_train.value_counts())
print("\n--- Churn Distribution in Test Set (y_test) ---")
print(y_test.value_counts())
```

 Numerical columns identified: ['SeniorCitizen', 'tenure', 'MonthlyCharges', 'TotalCharges']  
Categorical columns identified: ['gender', 'Partner', 'Dependents', 'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity']  
  
Training set shape: (5625, 19)  
Testing set shape: (1407, 19)  
  
--- Churn Distribution in Training Set (y\_train) ---  
Churn  
0 4130  
1 1495  
Name: churn, dtype: int64  
  
--- Churn Distribution in Test Set (y\_test) ---  
Churn  
0 1033  
1 374  
Name: churn, dtype: int64

4. Model Building With our data prepared, we'll now build and train various classification models. We'll use Logistic Regression (a simple linear model), Random Forest (an ensemble tree-based model), and XGBoost (a powerful gradient boosting model). Each model will be integrated into a pipeline with our preprocessor for seamless training.

```
# CODE BLOCK 4: Model Building

# --- Initialize Models ---
models = {
    'Logistic Regression': LogisticRegression(random_state=42, solver='liblinear', max_iter=1000), # max_iter for convergence
    'Random Forest': RandomForestClassifier(random_state=42),
    'XGBoost': XGBClassifier(random_state=42, use_label_encoder=False, eval_metric='logloss') # eval_metric for XGBoost future warning
}

# --- Create and Train Pipelines ---
trained_models = {}
for name, model in models.items():
    print(f"\n--- Training {name} ---")
    pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                               ('classifier', model)])
    pipeline.fit(X_train, y_train)
    trained_models[name] = pipeline
    print(f"{name} training complete. ✅")
```



```
--- Training Logistic Regression ---
Logistic Regression training complete. ✅

--- Training Random Forest ---
Random Forest training complete. ✅

--- Training XGBoost ---
XGBoost training complete. ✅
```

5. Model Evaluation After training, it's crucial to evaluate how well our models perform on unseen data. We'll use several key metrics for classification tasks, including Accuracy, Precision, Recall, F1-Score, and ROC-AUC. We'll also visualize the Confusion Matrix and ROC Curves, and analyze Feature Importance for tree-based models.

#### # CODE BLOCK 5: Model Evaluation

```
results = {}
for name, pipeline in trained_models.items():
    print(f"\n--- Evaluating {name} ---")
    y_pred = pipeline.predict(X_test)
    y_proba = pipeline.predict_proba(X_test)[:, 1] # Probability of churn (positive class)

    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    roc_auc = roc_auc_score(y_test, y_proba)

    results[name] = {
        'Accuracy': accuracy,
        'Precision': precision,
        'Recall': recall,
        'F1-Score': f1,
        'ROC-AUC': roc_auc
    }

    print(f"Accuracy: {accuracy:.4f} 🌟")
    print(f"Precision: {precision:.4f}")
    print(f"Recall: {recall:.4f}")
    print(f"F1-Score: {f1:.4f}")
    print(f"ROC-AUC: {roc_auc:.4f}")

    # --- Confusion Matrix ---
    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(6, 5))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                xticklabels=['No Churn', 'Churn'],
                yticklabels=['No Churn', 'Churn'])
    plt.title(f'Confusion Matrix for {name}')
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.show()

    # --- ROC Curve ---
    fpr, tpr, thresholds = roc_curve(y_test, y_proba)
    roc_auc_val = auc(fpr, tpr)
    plt.figure(figsize=(6, 5))
    plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc_val:.2f})')
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title(f'Receiver Operating Characteristic (ROC) Curve for {name}')
    plt.legend(loc="lower right")
    plt.show()

    # --- Feature Importance (for tree-based models like Random Forest, XGBoost) ---
    if hasattr(pipeline.named_steps['classifier'], 'feature_importances_'):
        importances = pipeline.named_steps['classifier'].feature_importances_
        # Get feature names after one-hot encoding
        ohe = pipeline.named_steps['preprocessor'].named_transformers_['cat'].named_steps['onehot']
        # Check if ohe successfully fitted, otherwise get_feature_names_out might fail on initial call
        try:
            categorical_feature_names = ohe.get_feature_names_out(categorical_cols)
        except AttributeError: # Fallback if get_feature_names_out isn't directly available or fitted
            categorical_feature_names = ohe.categories_[0] # simplified if ohe.categories_ is directly available for single category

        all_feature_names = numerical_cols + list(categorical_feature_names)

        feature_importance_df = pd.DataFrame({'feature': all_feature_names, 'importance': importances})
        feature_importance_df = feature_importance_df.sort_values(by='importance', ascending=False)

        plt.figure(figsize=(10, 6))
        sns.barplot(x='importance', y='feature', data=feature_importance_df.head(15), palette='viridis') # Top 15 features
        plt.title(f'Top 15 Feature Importance for {name} 🌟')
        plt.xlabel('Importance')
        plt.ylabel('Feature')
        plt.tight_layout()
        plt.show()

    elif hasattr(pipeline.named_steps['classifier'], 'coef_'): # For Logistic Regression
        coefficients = pipeline.named_steps['classifier'].coef_[0]
        # Get feature names after one-hot encoding
        ohe = pipeline.named_steps['preprocessor'].named_transformers_['cat'].named_steps['onehot']
        try:
            categorical_feature_names = ohe.get_feature_names_out(categorical_cols)
        except AttributeError:
```

```

categorical_feature_names = ohe.categories_[0]

all_feature_names = numerical_cols + list(categorical_feature_names)

coef_df = pd.DataFrame({'feature': all_feature_names, 'coefficient': coefficients})
coef_df['abs_coefficient'] = np.abs(coef_df['coefficient'])
coef_df = coef_df.sort_values(by='abs_coefficient', ascending=False)

plt.figure(figsize=(10, 6))
sns.barplot(x='coefficient', y='feature', data=coef_df.head(15), palette='coolwarm')
plt.title(f'Top 15 Feature Coefficients for {name} 💡')
plt.xlabel('Coefficient Value')
plt.ylabel('Feature')
plt.tight_layout()
plt.show()

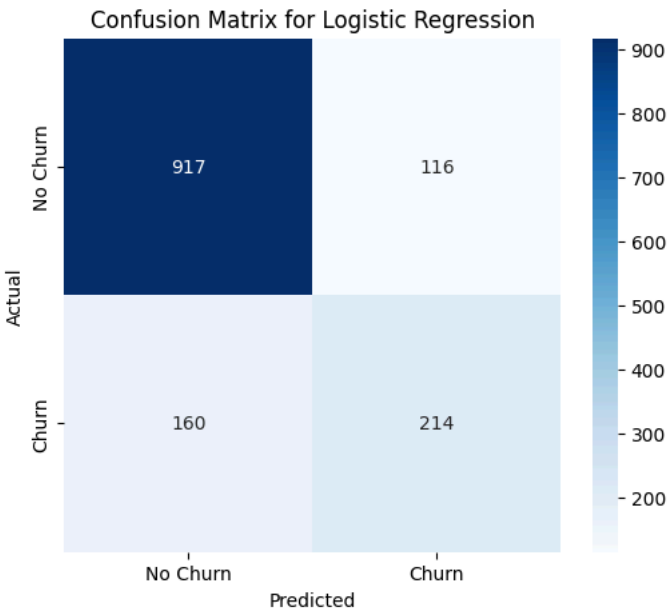
# --- Overall Model Comparison ---
print("\n--- Model Performance Comparison ---")
results_df = pd.DataFrame(results).T
print(results_df)

# You might want to choose the best model based on a specific metric (e.g., F1-Score or ROC-AUC for imbalanced churn data)
best_model_name = results_df['ROC-AUC'].idxmax()
print(f"\nBest model based on ROC-AUC: {best_model_name} 🏆")
best_pipeline = trained_models[best_model_name]

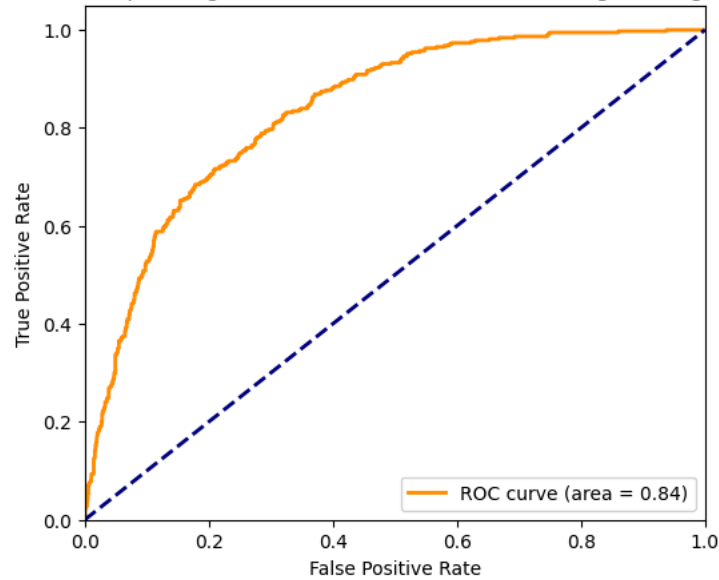
```



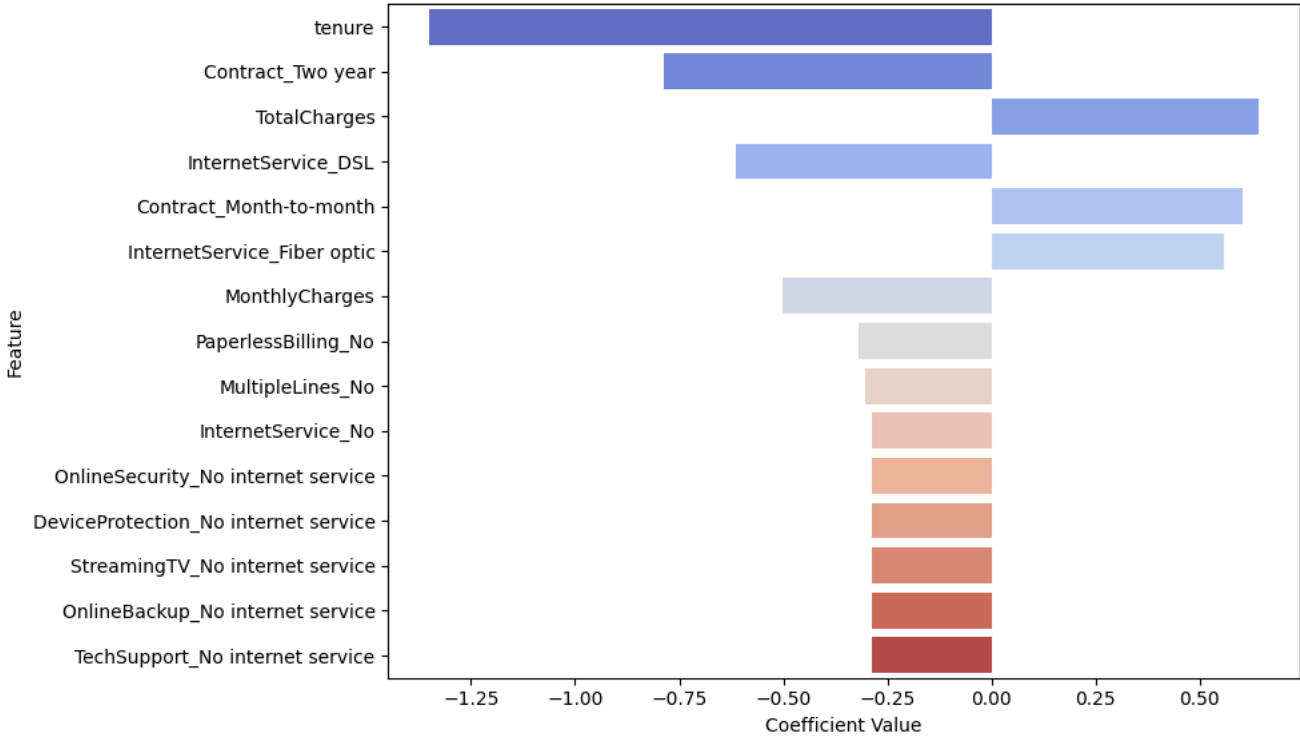
--- Evaluating Logistic Regression ---  
Accuracy: 0.8038 🌟  
Precision: 0.6485  
Recall: 0.5722  
F1-Score: 0.6080  
ROC-AUC: 0.8359



Receiver Operating Characteristic (ROC) Curve for Logistic Regression



Top 15 Feature Coefficients for Logistic Regression



--- Evaluating Random Forest ---  
Accuracy: 0.7875 🌟  
Precision: 0.6307  
Recall: 0.4840  
F1-Score: 0.5477  
ROC-AUC: 0.8137

